

Shouri Addepally
ECE 220 Honors
05/03/2020

MP 4 Honors: Report

Questions:

1. Below are a few benchmark results:

```
[shouria2@linux-a2 mp4h]$ python benchmark.py
Benchmark success!
Time taken for serial:      0.596627950668
Time taken for parallel:    0.597227096558
Speedup:                    1.0 X
[shouria2@linux-a2 mp4h]$ python benchmark.py
Benchmark success!
Time taken for serial:      0.6416451931
Time taken for parallel:    0.650228977203
Speedup:                    0.99 X
[shouria2@linux-a2 mp4h]$ python benchmark.py
Benchmark success!
Time taken for serial:      0.677536010742
Time taken for parallel:    0.595946073532
Speedup:                    1.14 X
[shouria2@linux-a2 mp4h]$ python benchmark.py
Benchmark success!
Time taken for serial:      0.675286054611
Time taken for parallel:    0.592781066895
Speedup:                    1.14 X
```

2. While generating my threading model, I decided to keep track of the number of levels that had already searched in my binary tree using `SWITCH_TO_SERIAL_FACTOR` which was set to 4. Until the search crossed a level of 4, my threading model would continue to spawn parallel threads and would use the parallelized searching approach. If this “threshold” of levels was crossed, I used serialization by calling the `recursive_helper` which would stop spawning threads and use serialization to finish searching the rest of the binary tree. The reason why I used this threading model is because the ‘`pthread_create`’ and ‘`pthread_join`’ operations are costly and as we progress down the tree, we will slow down in the searching process if we try to parallelize. As a result, a more logical approach is to parallelize the initial levels of the tree and if we get to a larger amount of levels in the binary tree we can perform serialization.
3. While I built my tree, I used a random number generator to fill in values on both the right and left branches of each node. While building my tree, I also used recursion to continue filling out random values on every right and left branch of the tree. As part of the tree generation, I also built an array which stored all the values that were generated as part of the tree. It is important to note that I set the rightmost leaf-node value with -1000, and this was the value I would search for in both the `find_serial` and `find_parallel` functions.
4. If our input tree looks like a linked list we cannot improve the runtime. This is because the binary tree does not have the invariant of being generated in any particular order. As a result, there are possibilities where we repeatedly traverse the linked list inefficiently which will cause a slowdown in the time taken to find a particular element.